



# Cache-Conscious Concurrent Hash Array Mapped Trie

Chi-Chun Chen  
Shaojie Wang  
Princeton Ferro




# Topics

- Hash Array-Mapped Tries
- Why Rust?
- Cache Consciousness
- Lock-free HAMT

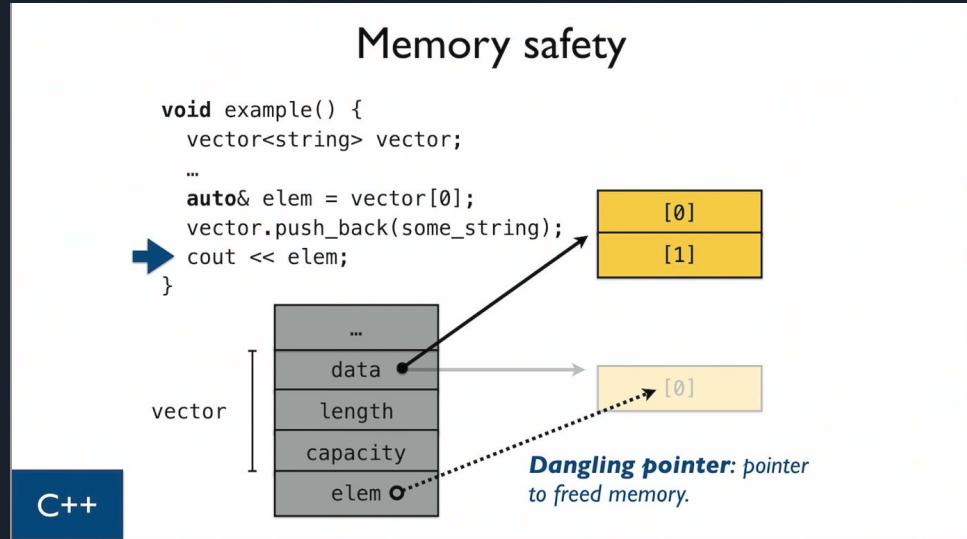
# Hash Array-Mapped Tries

1. A HAMT is a data structure where the data is stored at different levels.
2. A 64-bit hash of the key is divided into groups of 4 or 2 bits, and each group indexes into the next level.



# Hash Array-Mapped Tries

- Space efficient
- Persistent data structure. Get rid of dangling pointer issues.
- Small  $O(1)$  costs for principle operations (insertion, search, removal), and guaranteed small bounded worst case times. (Bagwell, 2001, Ideal Hash Trees.)





# Why Rust?

- Same niche as C/C++
  - No garbage collector
- Backed by llvm
- Static + Strong type
- Smart pointer by default
- Safe



# Why Rust?

- No concurrent Data Structure in standard library
- Few cache conscious implementation




# Cache-Consciousness


- Design a memory layout
  - Static Memory layout
    - Clustering

Root

insert  $\left( \frac{010}{2} \quad \frac{110}{6} \quad \frac{101}{5} , 100 \right)$




Layout :

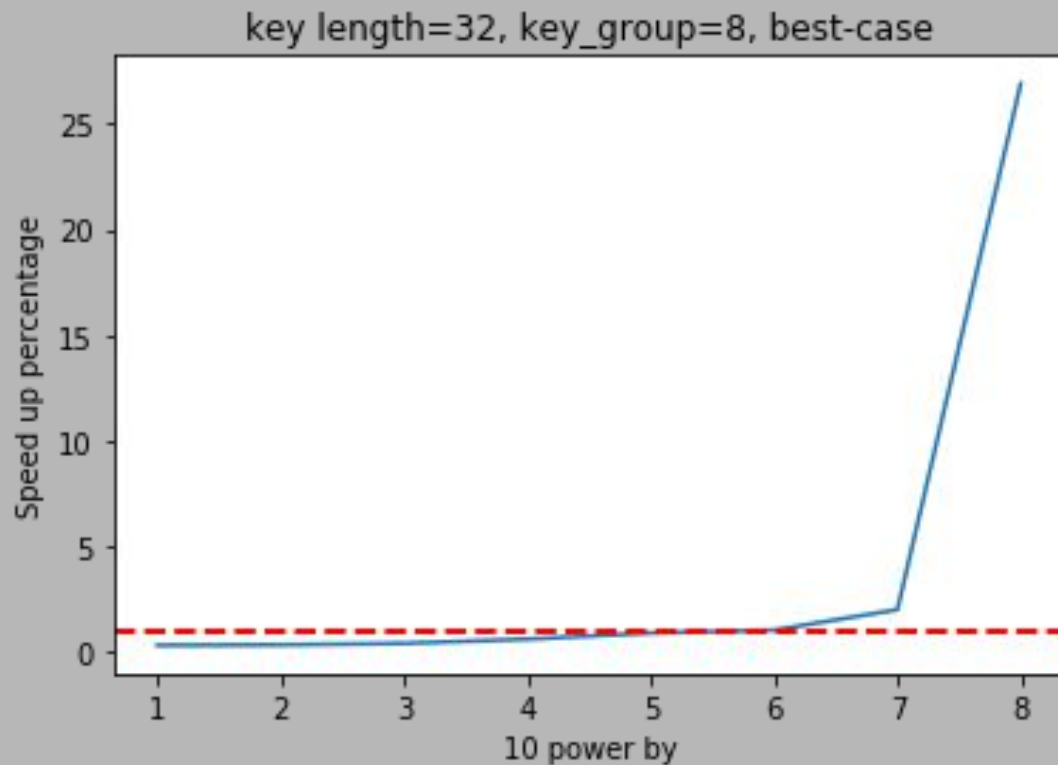




# Cache-conscious trie w/o concurrency



# Cache-conscious trie w/o concurrency





# Cache-Consciousness

- Design a memory layout
  - Static Memory layout
    - Clustering
- Dynamic Memory layout
  - Packing the data in the order of "First Appearance"
  - Group packing
  - Reference: Chen Ding, Ken Kennedy: Improving Cache Performance in Dynamic Application through Data and Computation Reorganization at Run time



# Cache-Consciousness


- Design a memory layout
  - Static Memory layout
    - Clustering
- Dynamic Memory layout
  - Packing the data in the order of "First Appearance"
  - Group packing
  - Reference: Chen Ding, Ken Kennedy: Improving Cache Performance in Dynamic Application through Data and Computation Reorganization at Run time

# Lock-free HAMT

## Nodes are immutable

Each node has an associated transaction field that indicates there is a concurrent modification request.

When another thread sees this, it helps complete the transaction.





# Lock-free HAMT: Cache and Memory Layout

## Cache Trie

We also cache the largest level of the trie.  
After a certain number of misses,  
periodically, we change this level.


## Layout

We implement consecutive packing for Node  
elements, in the order of inserts.

# Results - Lockfree

**Tested against Rust's Concurrent HashMap.**

Though we are beaten by chashmap in most cases, our performance scales well, particularly when the number of threads exceeds the number of hardware contexts.



# Conclusion

In this project, we implemented concurrent cache-conscious hash array mapped trie. We first implement the memory layout for cache-consciousness, which gives the hash trie a speed up of 2500% in best case with 10,000,000 nodes. For the blocking version of hash trie, we can see a good scalability when node number goes larger. For lock-free version, we implemented a concurrent HAMT with a cache-trie for  $O(1)$  lookup in the best-case. Our implementations scale well with the number of threads.

Multi-thread programming:

